

CLASSIFICATION: THE NEAREST-NEIGHBORS ALGORITHM (K-NN)

AP

FROM THE INTRODUCTION:

1. Classification and class probability

Instance:

- a collection (dataset) of datapoints from \mathbf{X}
- a classification system $C = \{c_1, c_2, \dots, c_r\}$

Solution: classification function $\gamma : \mathbf{X} \rightarrow C$

Measure: misclassification

BINARY CLASSIFICATION

$r = 2$: positive and negative.

Misclassification is described by the *confusion matrix*, which scores the result of classification against labeled examples.

BINARY CLASSIFICATION IN 2D

With just two numerical dimensions, datapoint similarity can be interpreted as simple Euclidean distance.

Being very close \iff being very similar

Q: are 4 and 6 more similar to each other than -1 and +1?

Assumption: small changes in the values won't alter the classification, close points will receive the same classification


if a point is in close distance to a labeled one then assign the same class


THE NEAREST NEIGH. ALGORITHM

Take a set of labeled points (the examples), all others are *blank* at the moment.

Whenever a blank point has a nearest neighbor datapoint which is labeled, give it the same label

This is the NN, or 1-NN algorithm.





$$\gamma(\mathbf{x}) = y_i, i = \operatorname{argmin}_j ||\mathbf{x}_j - \mathbf{x}||$$


FROM 1-NN TO K-NN

Consider the k nearest neighbors


Assign the class that is the most common among them

Variation: consider each label relative to the effective distance of the neighbor.

1-NN



3-NN



LEARNING


Given the labeled examples, k-NN determines the areas around each example which give a certain class.

k-NN learns an area or *surface* and applies it in classification


A larger k does not always mean a better classification

INFLUENCE OF K


n_neighbors=1




n_neighbors=5



n_neighbors=10



n_neighbors=30



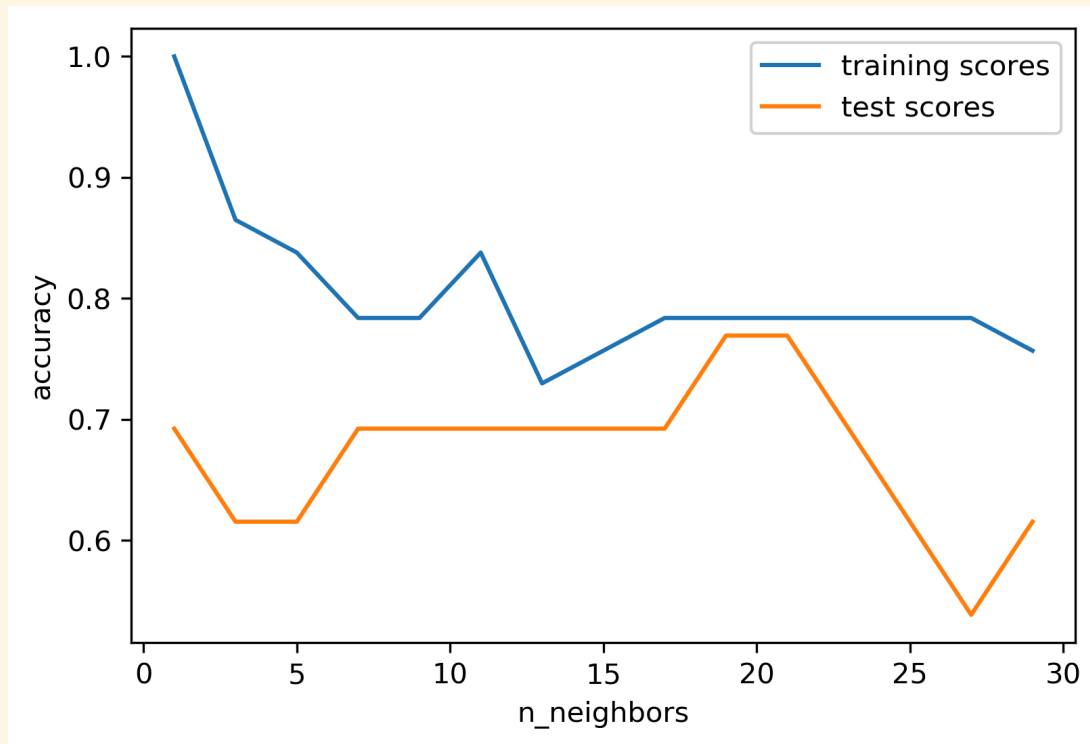
OBSERVATIONS

k-NN

- introduces us to *voting systems*
- is effective when the two classes are balanced, i.e., not *skewed*, in the dataset
- there is no standard way to choose k , yet it may greatly influence the outcome:
 - we face hyperparameter optimization.
- on large training data, even 1-NN approaches the *irreducible_error_rate* (2x).

TRADE-OFFS

Sometimes improving accuracy on the training data does not translate into improved accuracy in testing against *unseen* data



1-NN is perfect on training but 0.7 on test.

Increasing k does not improve much and *overfitting* creeps in.