LATENT DIMENSIONS

AP

DIMENSIONALITY REDUCTION

Problem 8: Data reduction

Instance:

- a collection (dataset) $\mathbf{D}$ of datapoints from $\mathbf{X}$, e.g., $\mathbb{R}^m$
- [a distinct independent variable x_i]

Solution: a projection of $\mathbf{D}$ onto $\mathbb{R}^n, n < m$

Measure: error in the estimation of x_i

Example: genre identification in consumer behaviour analysis

EXAMPLE

Olivetti faces are 64x64 binary matrices.

Through SVD we discovered that most singular values are in fact 0 or very small

by considering only the top 20 or so singular values we can obtain a very similar image with less data:

- A contains 64x64=4096 values
- $U\Sigma V^T$ contains 20x20 + 20 + 20x20 = 820 values (5:1 compression ratio)

WHY REDUCE DIMENSIONS?

- easier to store, quicker to process
- interpretation and visualisation
- remove redundant or noisy features
- escape the curse of dimensionality and go back to intuitive distance features
- discover hidden correlations/topics

WHEN REDUCE DIMENSIONS?

When any of the goals becomes important

Empirically: when we believe that data essentially represents the mixing of a smaller set of *feature* which are the real *axes* of the data.

There really are only two *features*: sci-fi and romance

	Matrix	Alien	Star Wars	Casablanca	Titanic
Joe	1	1	1	0	0
Jim	3	3	3	0	0
John	4	4	4	0	0
Jack	5	5	5	0	0
Jill	0	0	0	4	4
Jenny	0	0	0	5	5
Jane	0	0	0	2	2

Figure 11.6: Ratings of movies by users

HOW TO REDUCE DIMENSIONS?

High-level view:

- data lies on or near a low-dimensional space:
- axes of that space are effective representations of data

EXAMPLES: $D \longrightarrow d$

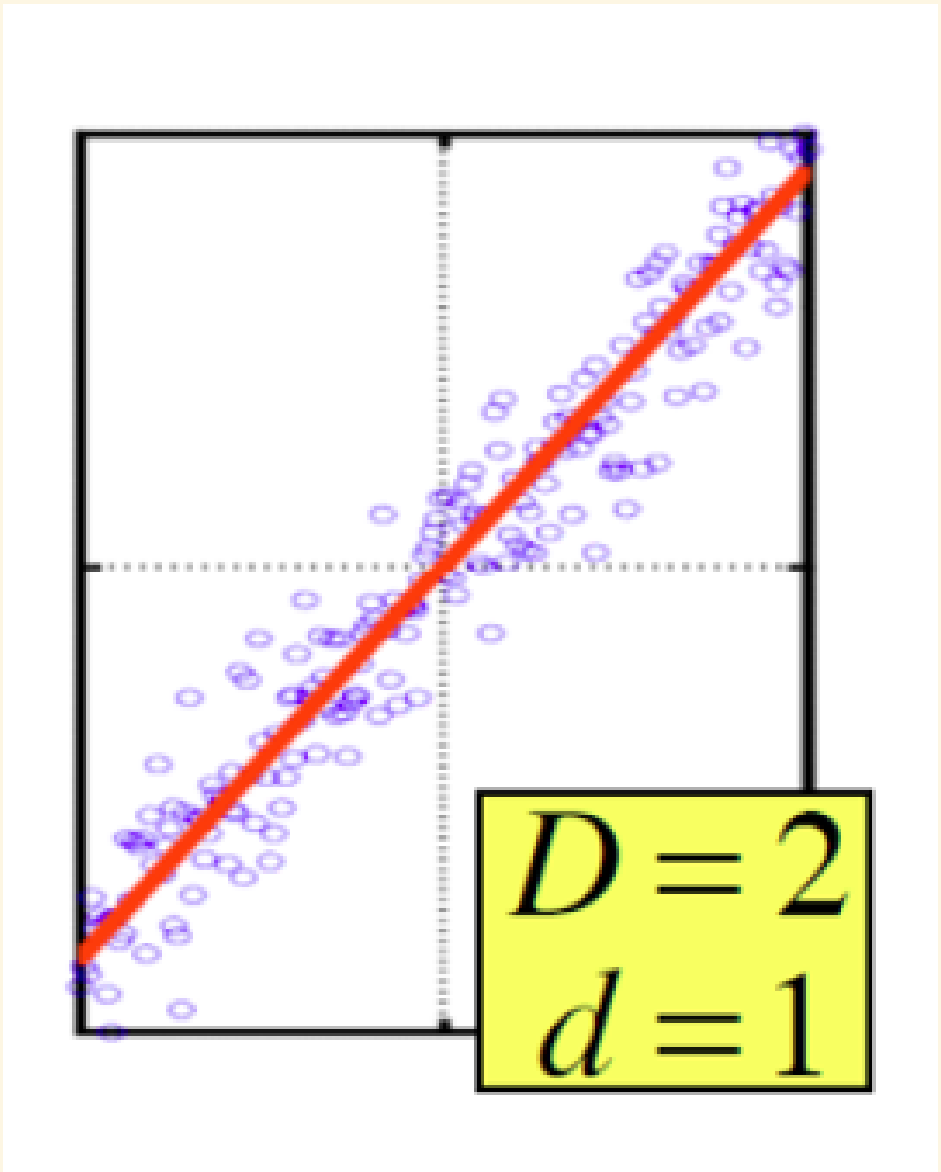

The data axes (in red) are almost never the original measurement axes

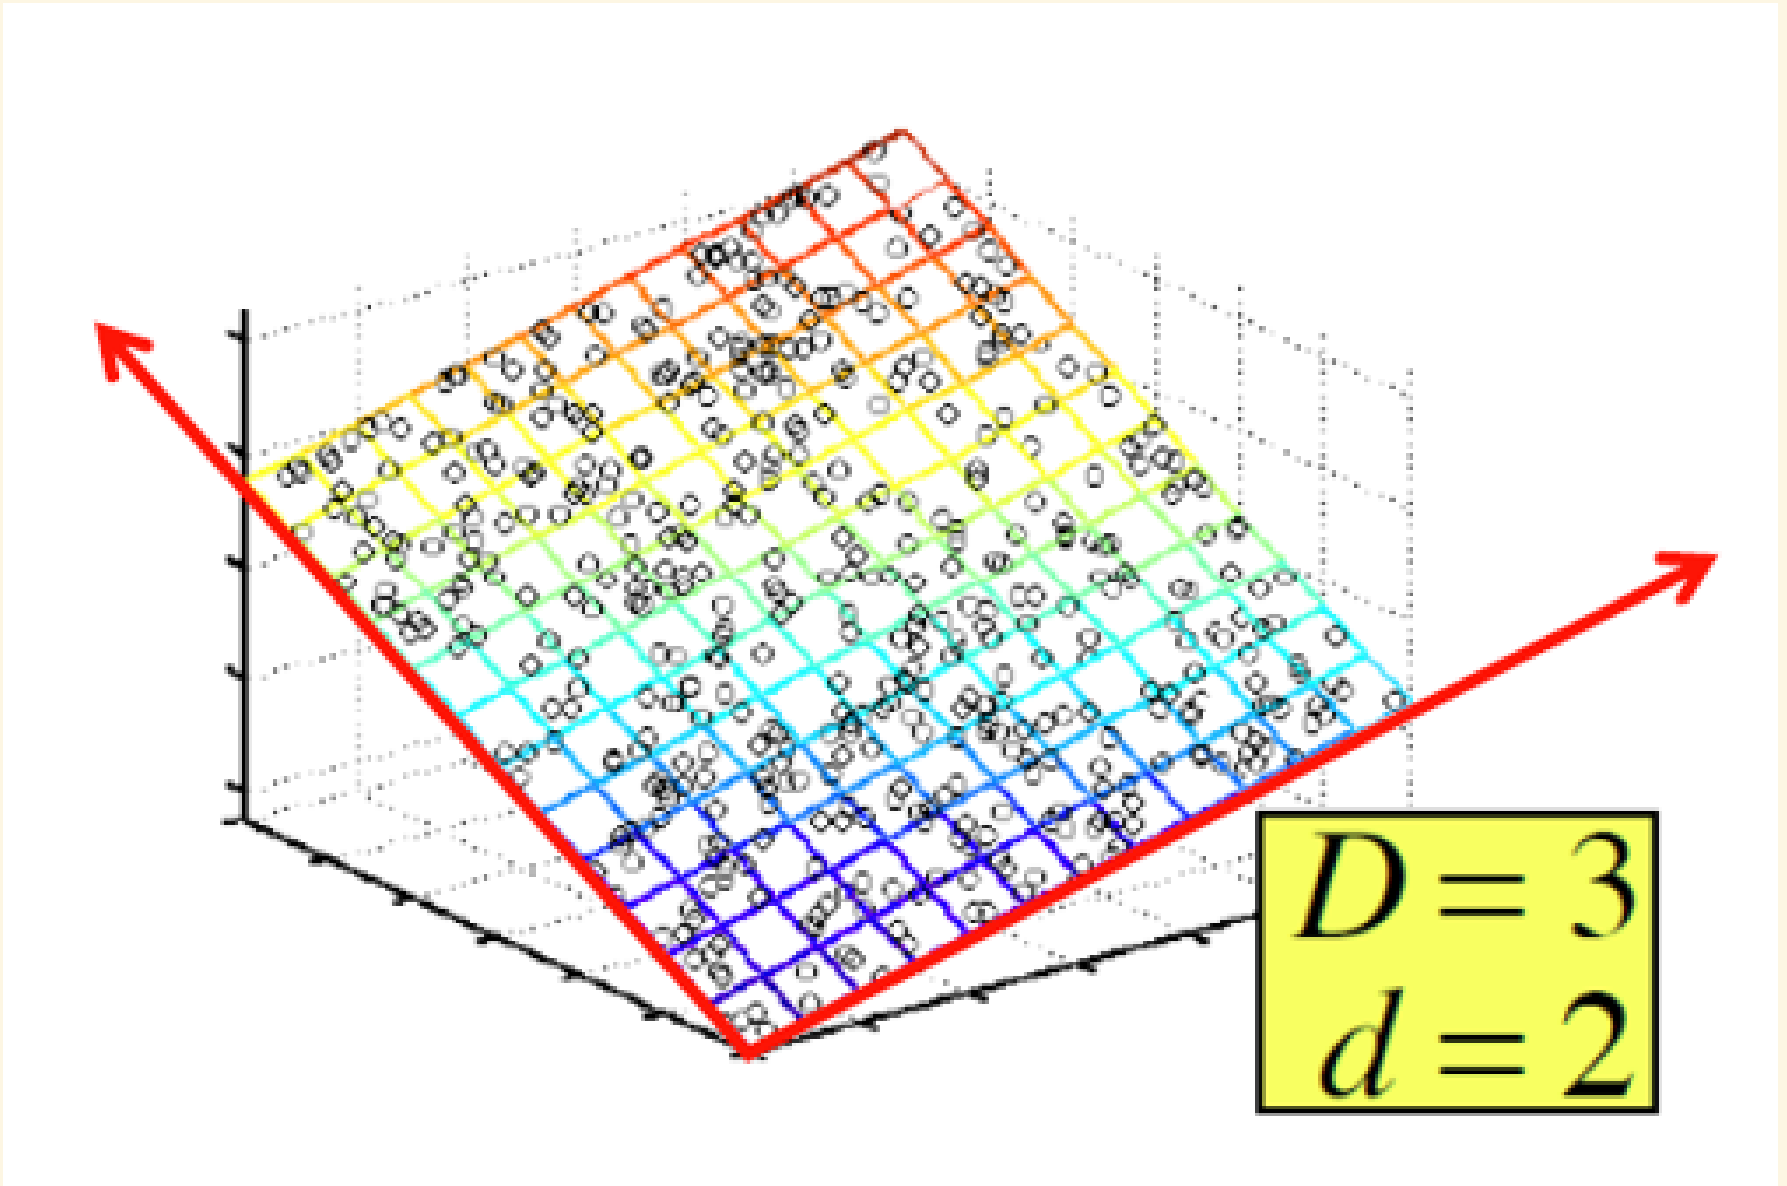

BANKGROUND: MATRIX RANK

Rank is an important feature/descriptor of a data matrix

Rank is the maximum number or colums (or rows) that are linearly independent.

Such independent cols/rows are *candidates for the new, reduced reference system* (the red axes)

Rectangular matrices with SVD: $r \leq \min\{n, m\}$

We can map data points is a completely-new, dataset-dependent representation!

The dataset-depent representation will be compact

Let's create a new, abstract *feature space*

EXAMPLE

Handmade dimensionality reduction, from the MMDS textbook

$$A_{3\times 3} = \begin{pmatrix} 1 & 2 & 1 \\ -2 & -3 & 1 \\ 3 & 5 & 0 \end{pmatrix}$$

$r = 2$ as the third row can be expressed as the first minus the second:

$$\begin{array}{rrrr} 1 & 2 & 1 & - \\ -2 & -3 & 1 & = \\ \hline \\ 3 & 5 & 0 & \end{array}$$

We create a new 2-d space where the axes are the first two rows:

$[1, 0, 0], [0, 1, 0], [0, 0, 1] \longrightarrow [1, 2, 1], [-2, -3, 1]$

The new rows are $[1, 0]$ $[0, 1]$ and $[1, 1]$

$$A'_{3\times 2} = \begin{pmatrix} 1 & 0 \\ 0 & 1 \\ 1 & -1 \end{pmatrix}$$

The new points work as selectors of the new axes: it's easy to go back from this space to the original, no loss of information/precision

DIMENSIONALITY REDUCTION

In real dataset mapping to a lower-dimensional space may introduce errors

in the $2 \rightarrow 1$ example below, instead of points we just take the measure (distance from the origin) of their projection on the red axis

The red axis is chosen as to minimise the error introduced by the $2 \rightarrow 1$ *projection*.

Data mining studies how to find such axes, called concepts

They capture some alignment which is inherent to the data.

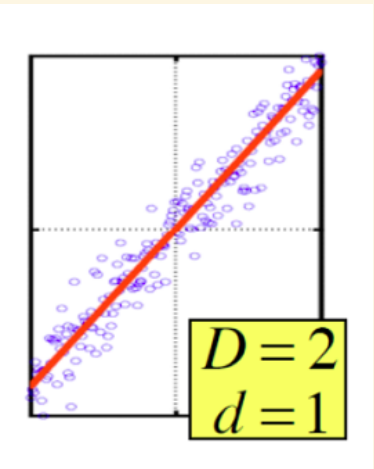

SVD

DECOMPOSITION

$$A_{(m\times n)} = U_{(m\times m)} D_{(m\times n)} V^T_{(n\times n)}$$

- U is a orthogonal m. of *left-singular* (col.) vectors
- D is a diagonal matrix of *singular values*
- V is a orthogonal m. of *right-singular* (col.) vectors

Suppose only r ($r < \min\{m, n\}$) singular values are non-zero

We can rewrite the decomposition as follows:

$$A_{(m\times n)} \approx U_{(m\times r)} D_{(r\times r)} V^T_{(r\times n)}$$

Suppose $r = 2$, visualise

$$A_{(m \times n)} \approx U_{(m \times 2)} D_{(2 \times 2)} V^T_{(2 \times n)}$$

$$A_{(m \times n)} \approx \sum_i \sigma_i U_i \circ V_i^T$$

where U_i is the i-th column of U and $\circ$ is matrix multiplication

Now A is represented as the sum of independent *factors* that were not explicit in the original data

EXAMPLE: USERS TO FILMS BECOMES USERS TO CONCEPTS TO FILMS

(sligthy different data)

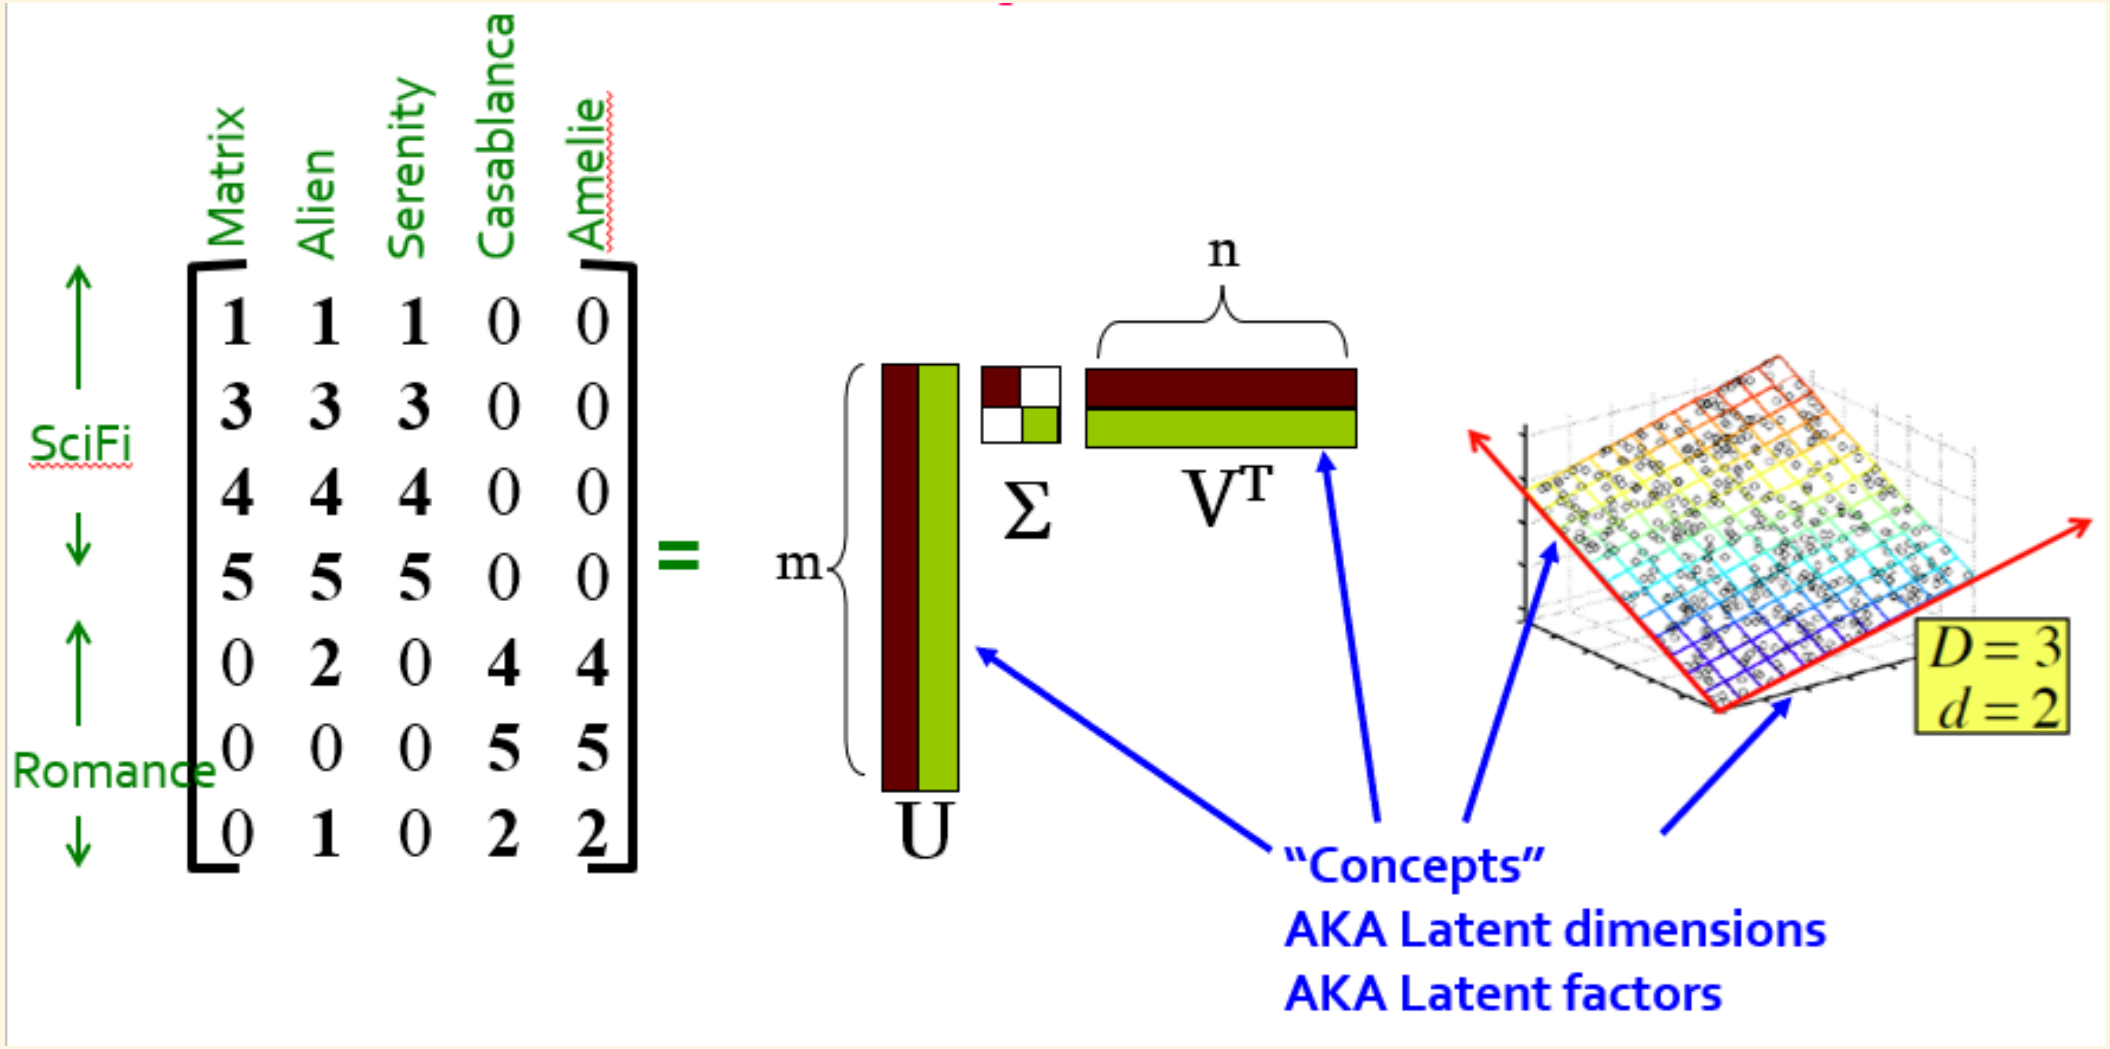

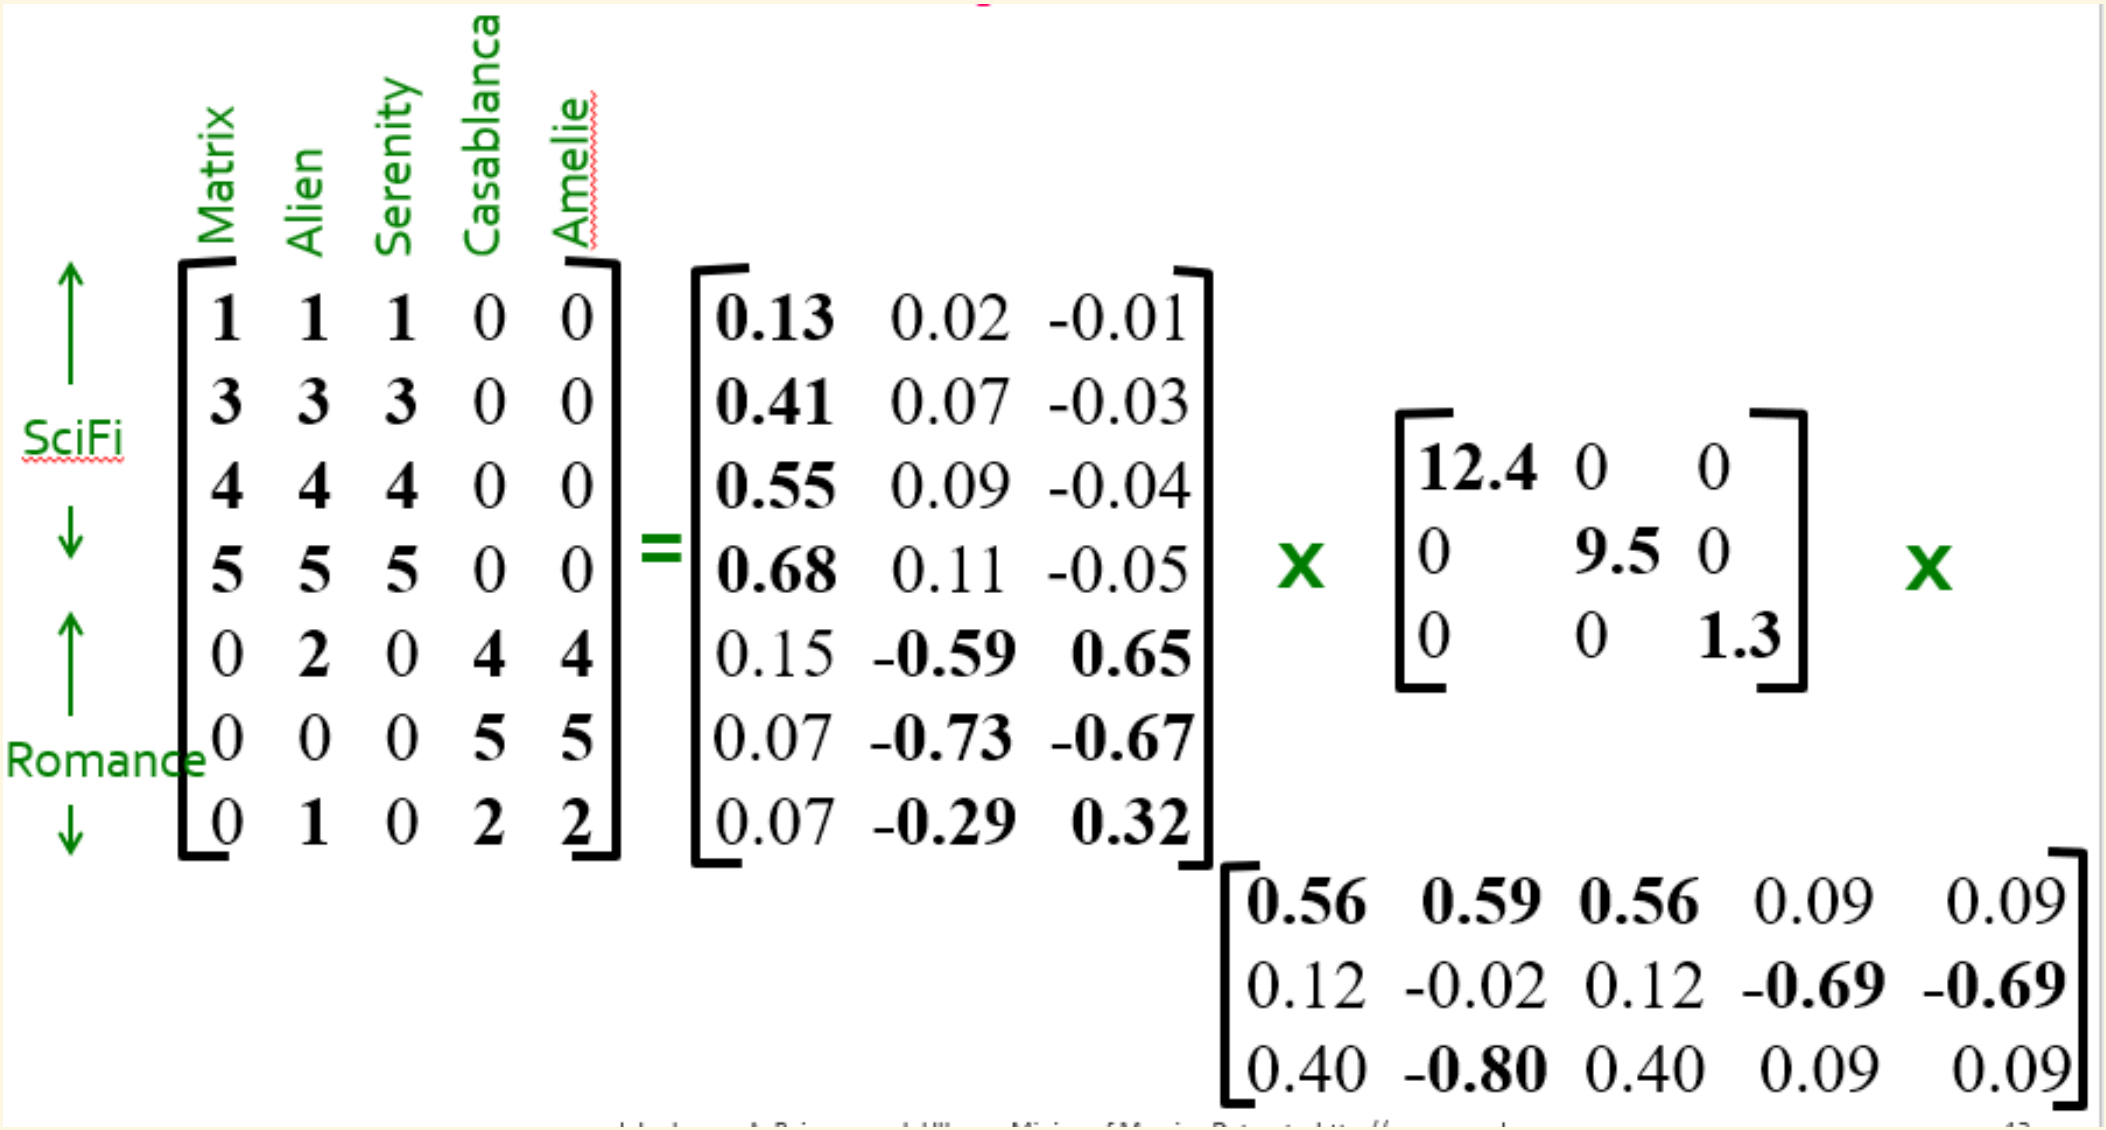
Matrix
Alien
Serenity
Casablanca
Amelie
SciFi
Romance
1 1 1 0 0
3 3 3 0 0
4 4 4 0 0
5 5 5 0 0
0 2 0 4 4
0 0 0 5 5
0 1 0 2 2
=
0.13 0.02 -0.01
0.41 0.07 -0.03
0.55 0.09 -0.04
0.68 0.11 -0.05
0.15 -0.59 0.65
0.07 -0.73 -0.67
0.07 -0.29 0.32
x
12.4 0 0
0 9.5 0
0 0 1.3
x
0.56 0.59 0.56 0.09 0.09
0.12 -0.02 0.12 -0.69 -0.69
0.40 -0.80 0.40 0.09 0.09

CONCEPTS

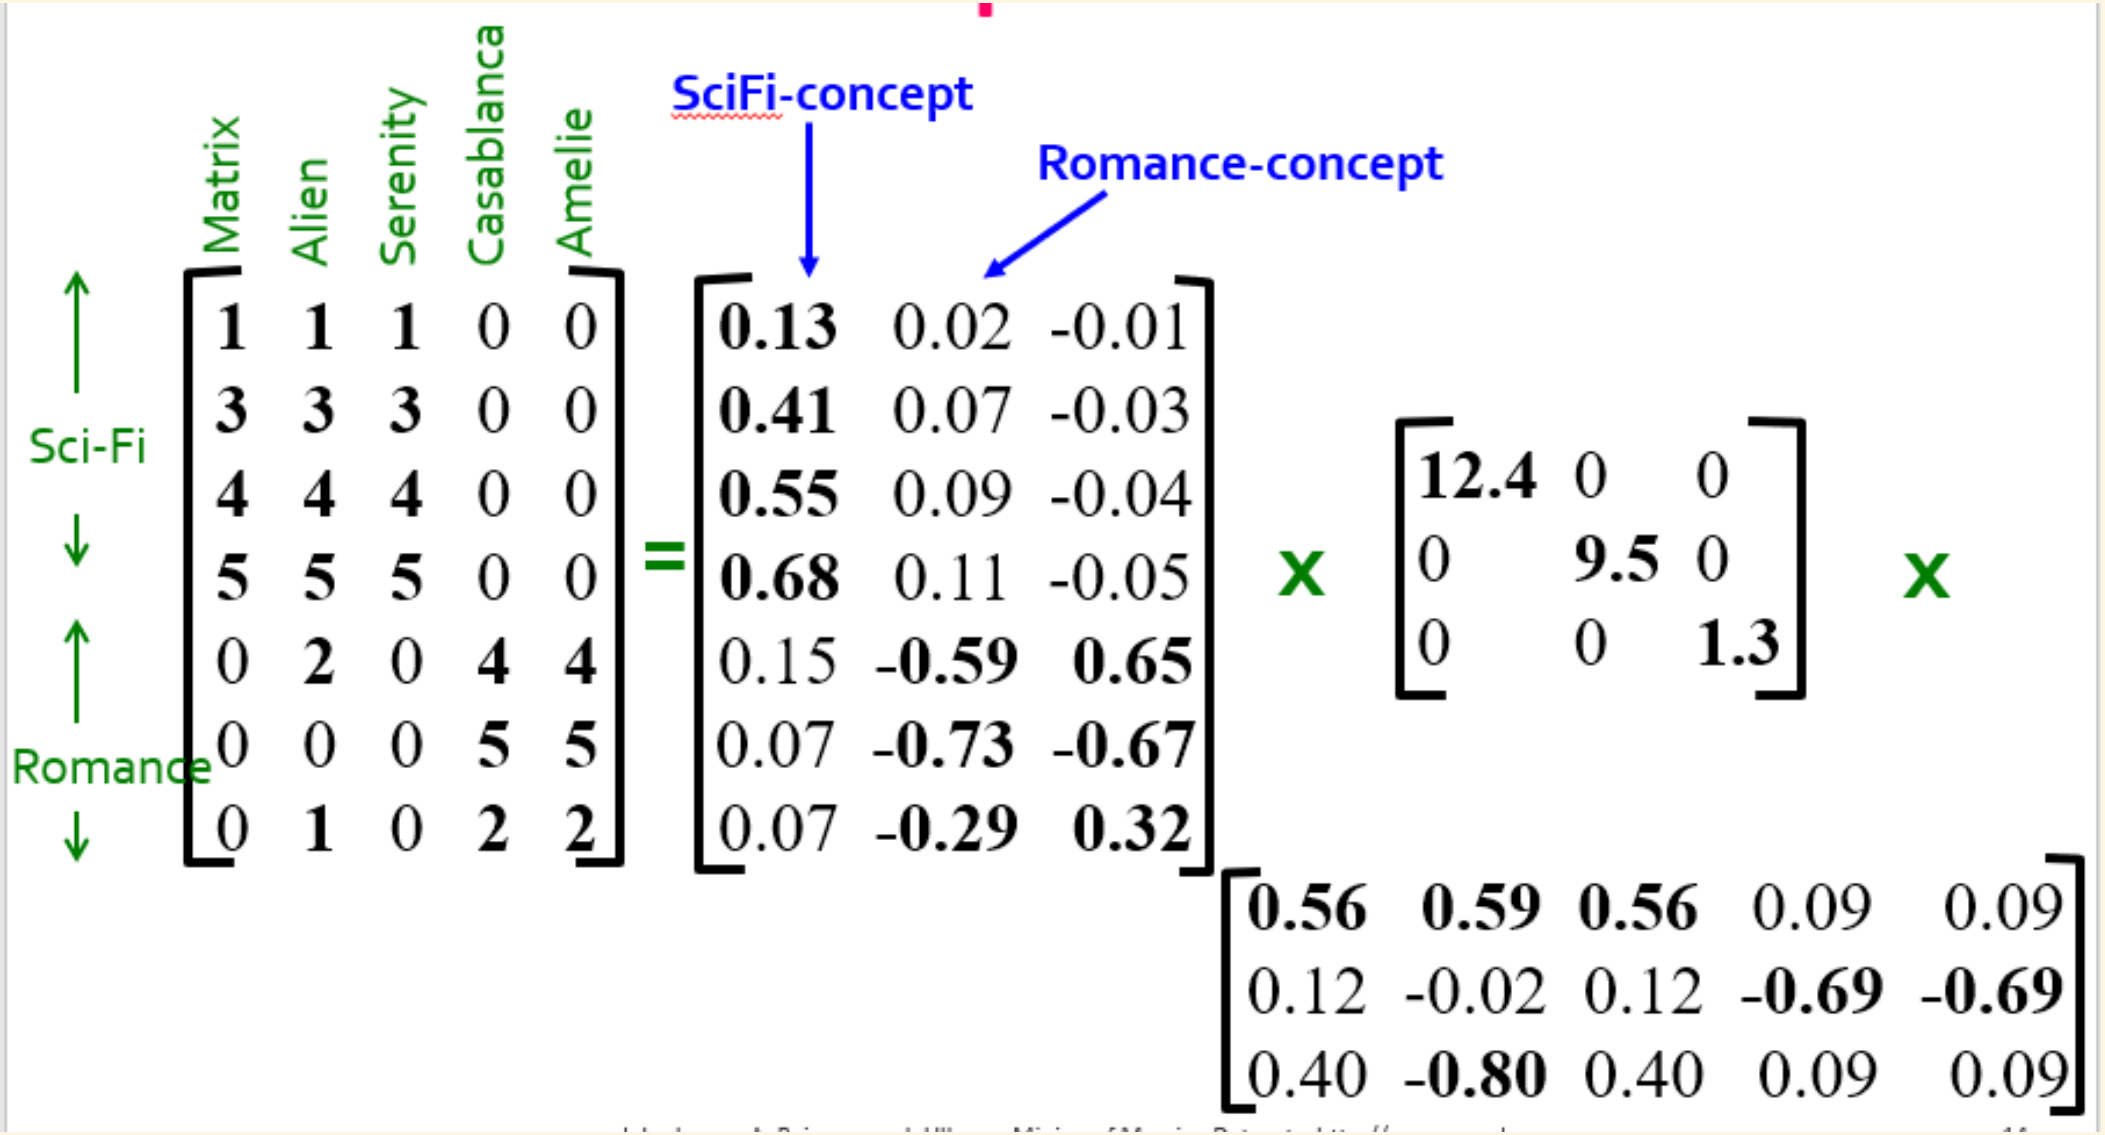
SciFi-concept
Romance-concept
Matrix
Alien
Serenity
Casablanca
Amelie
Sci-Fi
Romance
1 1 1 0 0
3 3 3 0 0
4 4 4 0 0
5 5 5 0 0
0 2 0 4 4
0 0 0 5 5
0 1 0 2 2
=
0.13 0.02 -0.01
0.41 0.07 -0.03
0.55 0.09 -0.04
0.68 0.11 -0.05
0.15 -0.59 0.65
0.07 -0.73 -0.67
0.07 -0.29 0.32
x
12.4 0 0
0 9.5 0
0 0 1.3
x
0.56 0.59 0.56 0.09 0.09
0.12 -0.02 0.12 -0.69 -0.69
0.40 -0.80 0.40 0.09 0.09

U is a user-to-concept similarity matrix

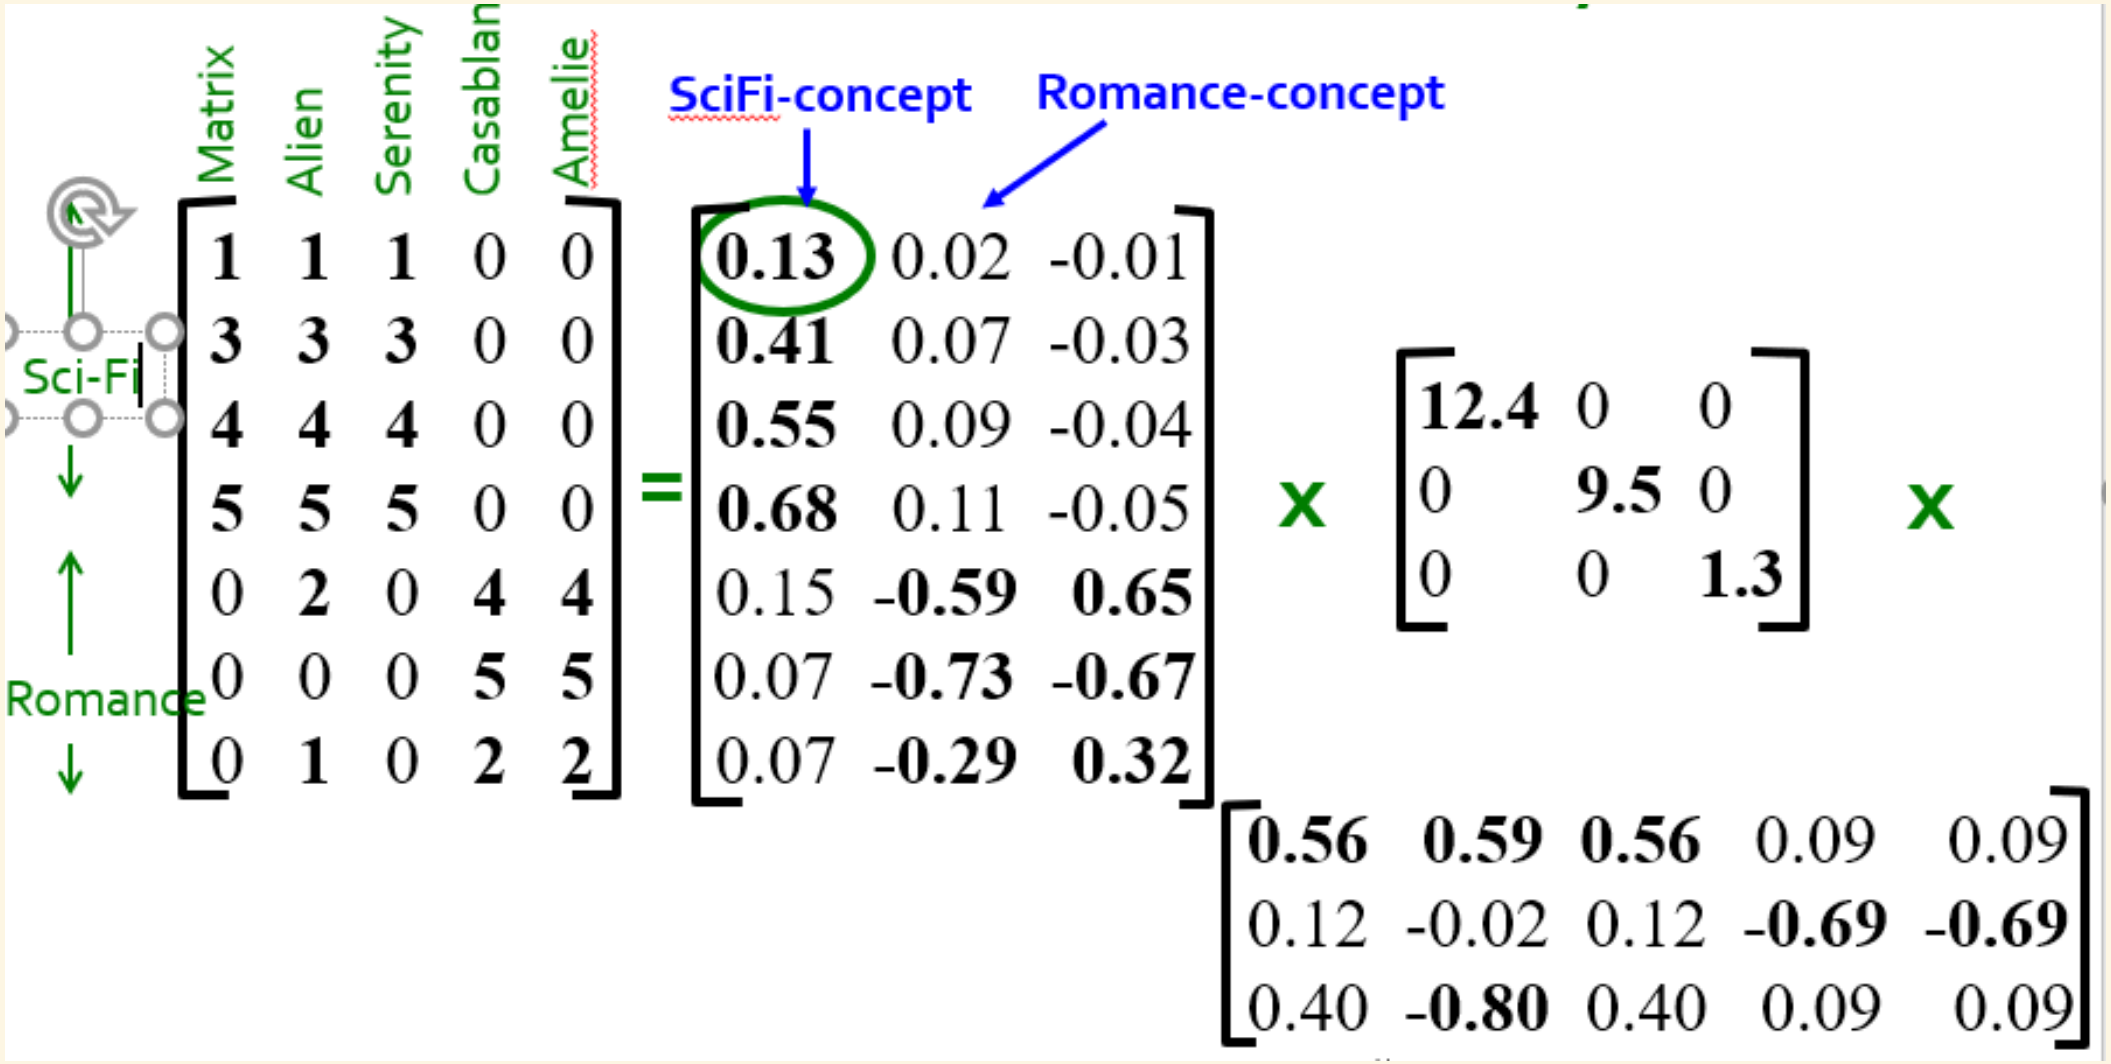

σs reveal the strength of each concept

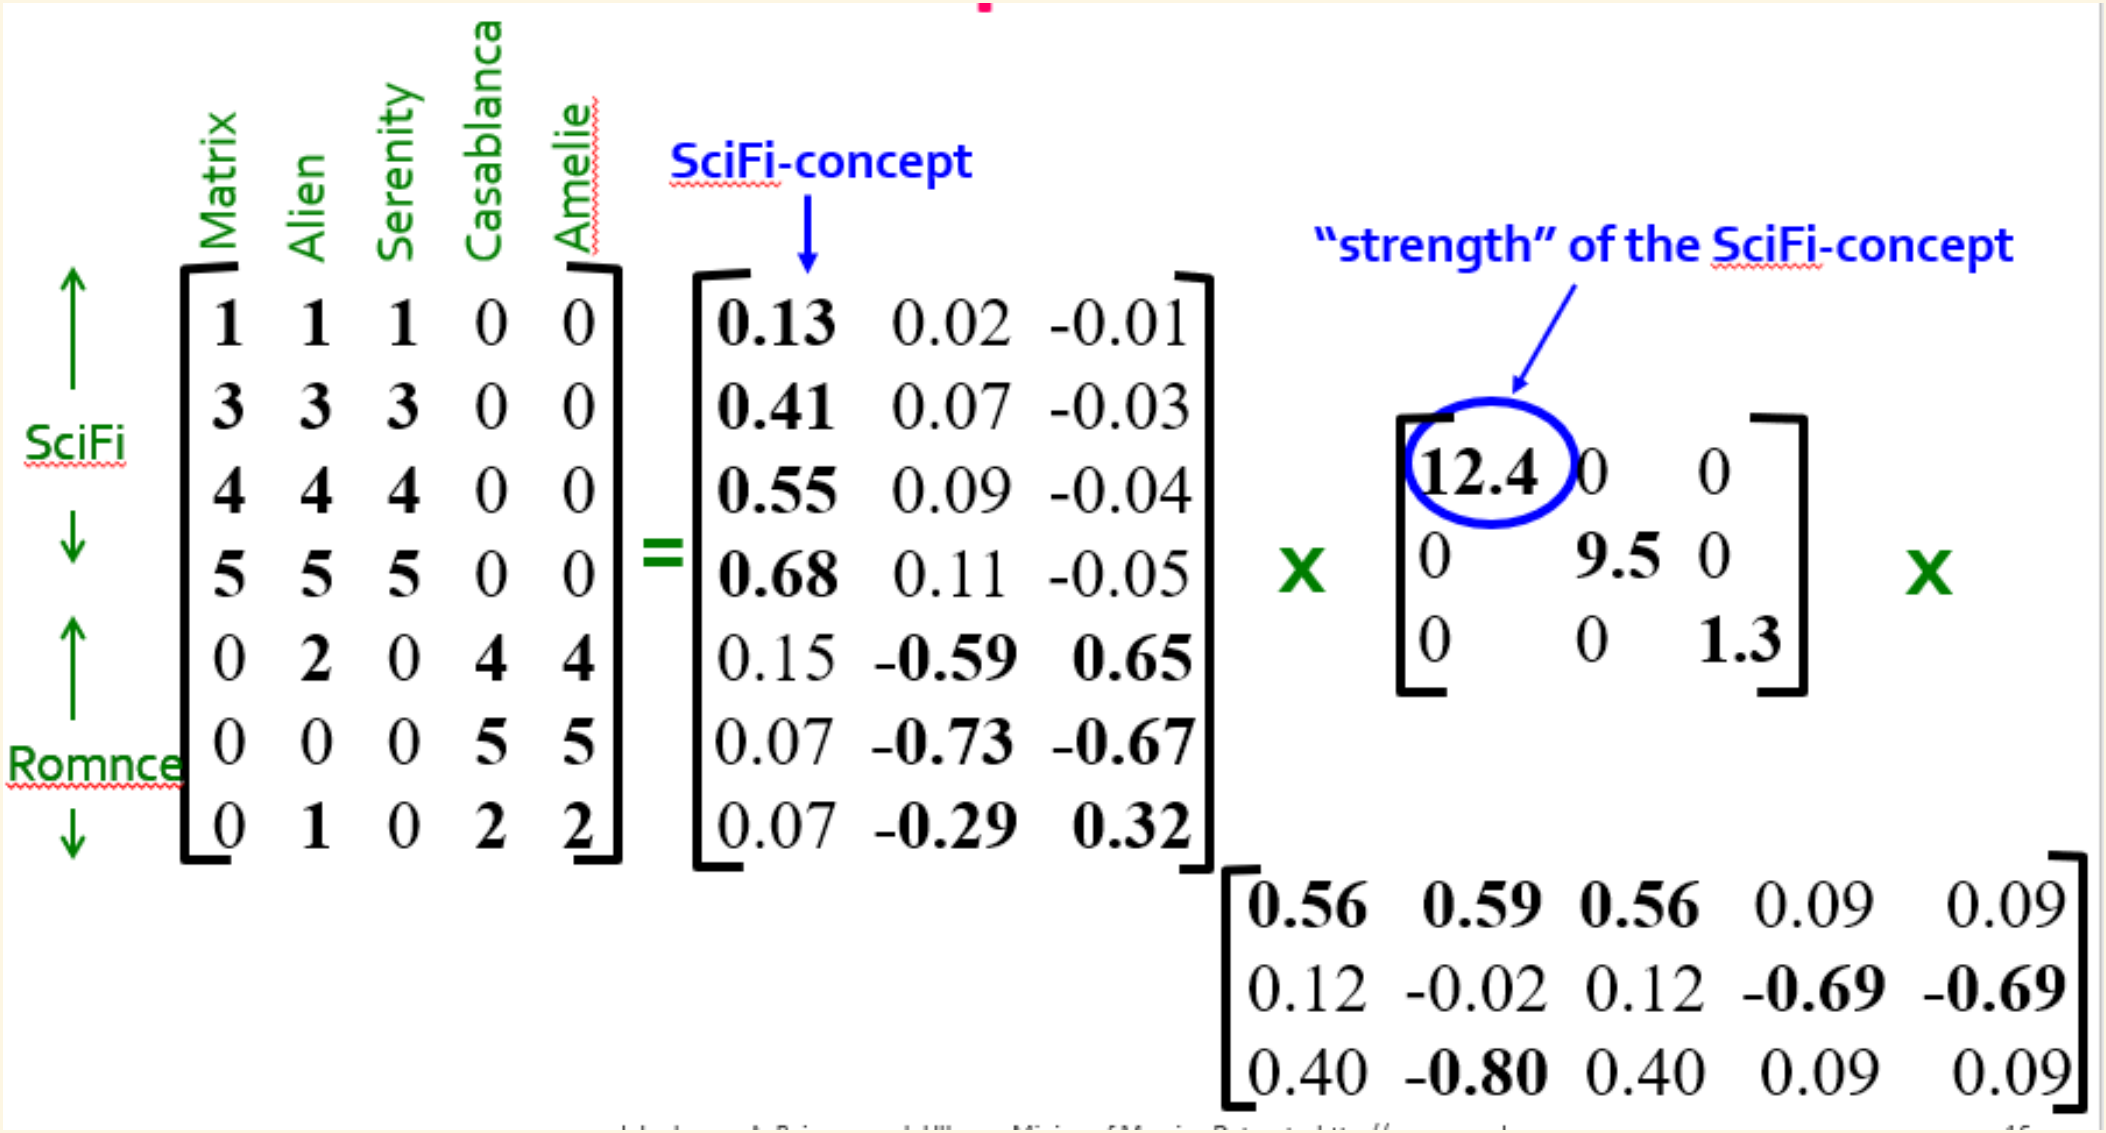

V^T is a concept-to-film similarity matrix

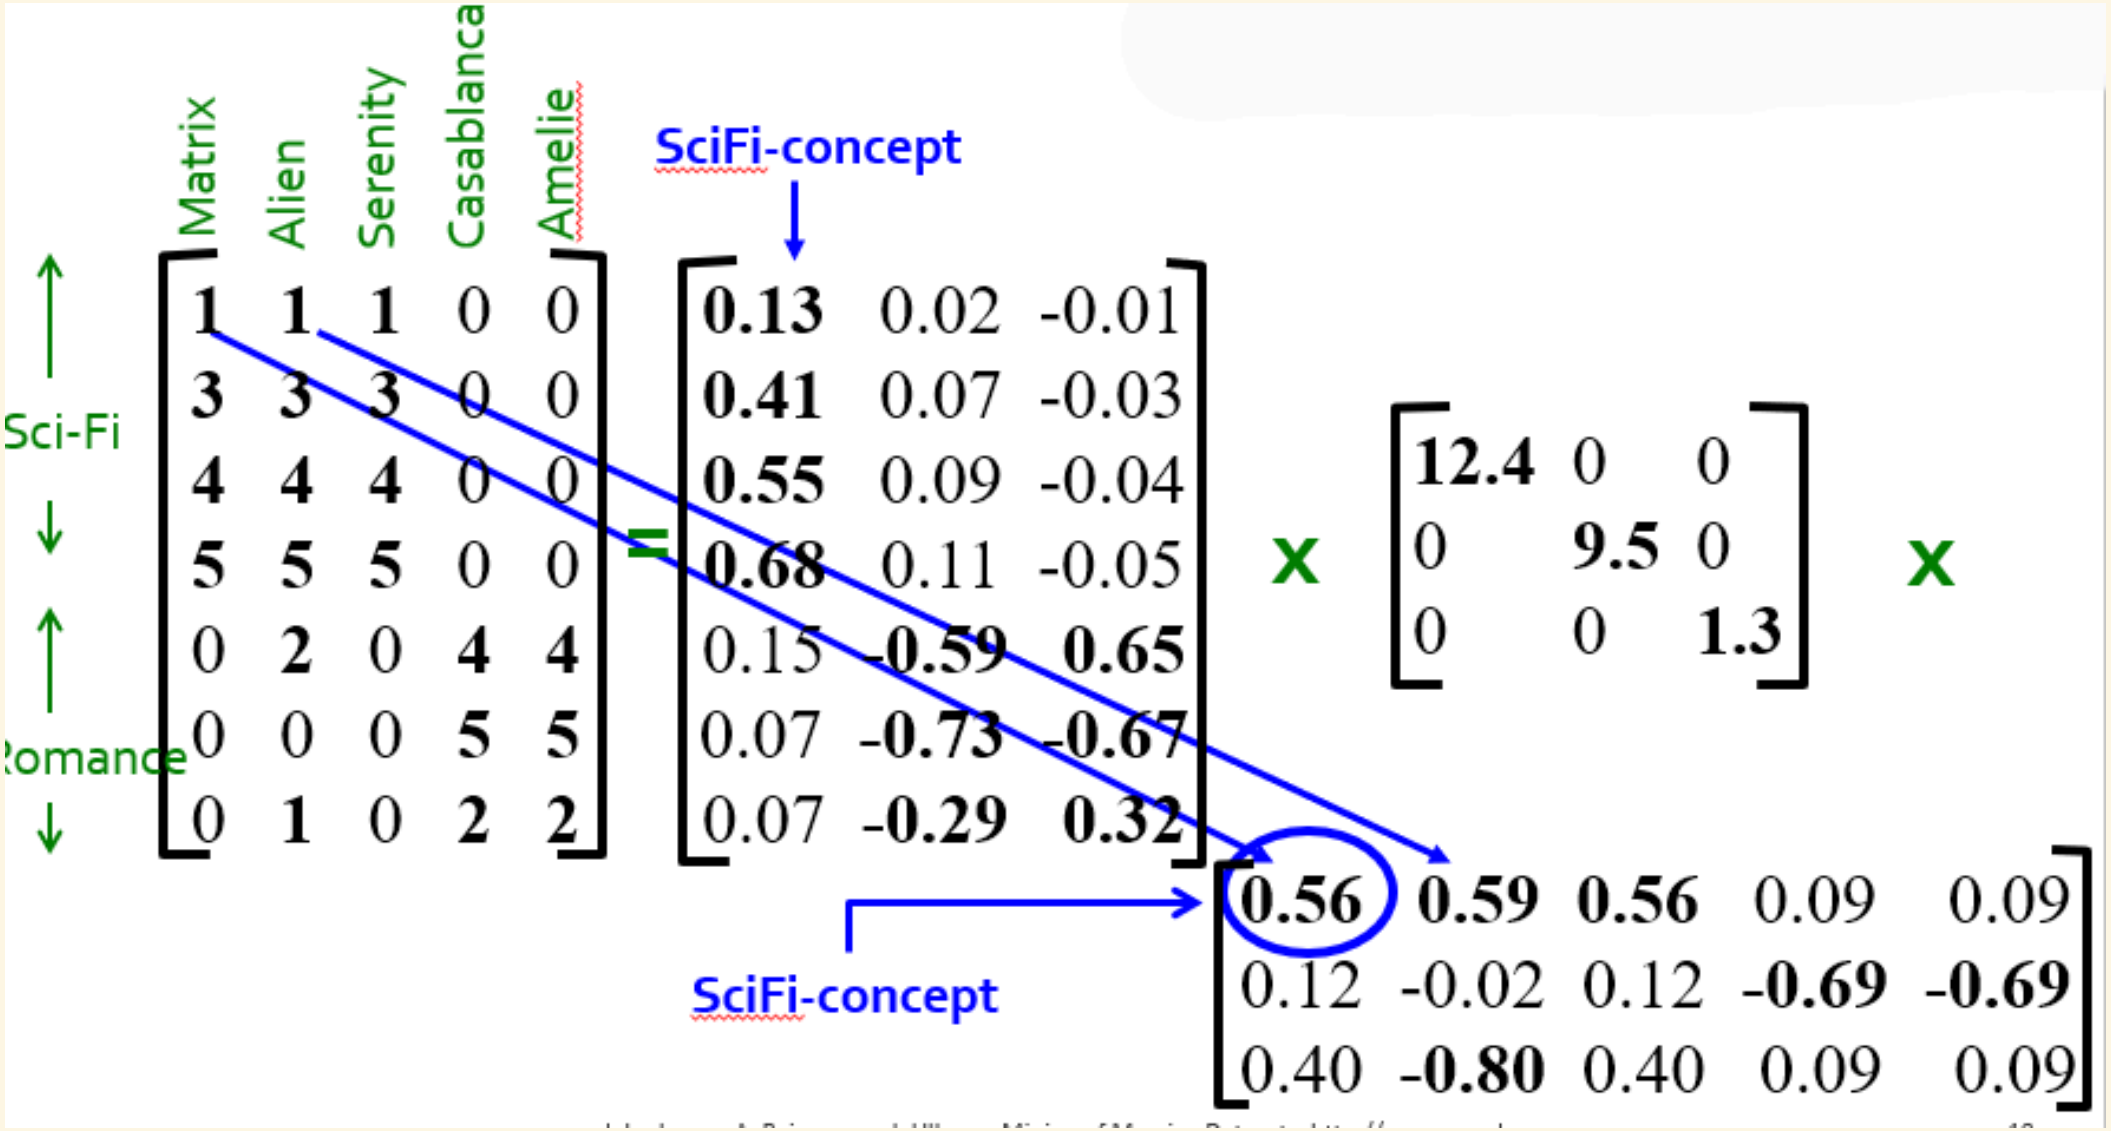

SVD INTERPRETATION -1

FIRST INTERPRETATION

The *singular vectors* that make up V (and U) are the new axes for projection

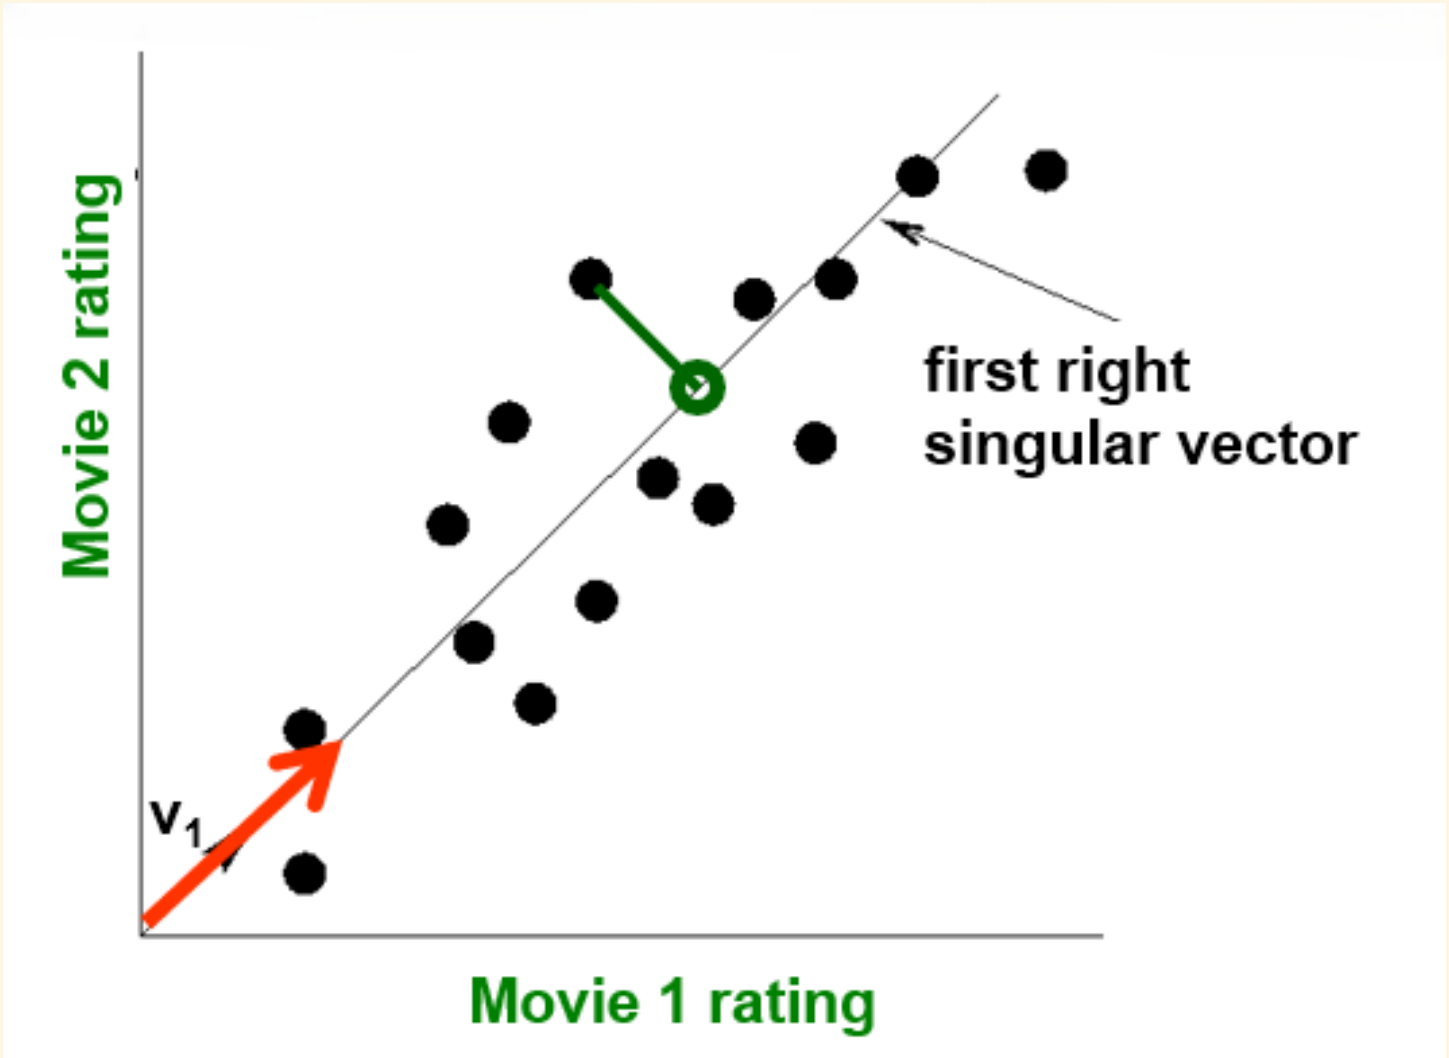

The *singular vectors* that make up V (and U) are the new axes for projection

They will minimise the *reconstruction error* (z is the the value obtained by the reduced SVD)

$$\epsilon = \sum_{i=1}^{m} \sum_{i=1}^{n} \|a_{ij} - z_{ij}\|$$

The sum of the green, orthogonal lines connecting the point to its compressed representation on the new axis is minimal

SVD INTERPRETATION -2

SECOND INTERPRETATION

Again, we use the singular vectors as the new axes

The σs indicate the *spread* (variance) of the projected points on the new axis

the more spread apart points are, the easier it will be to classify/find cluster them

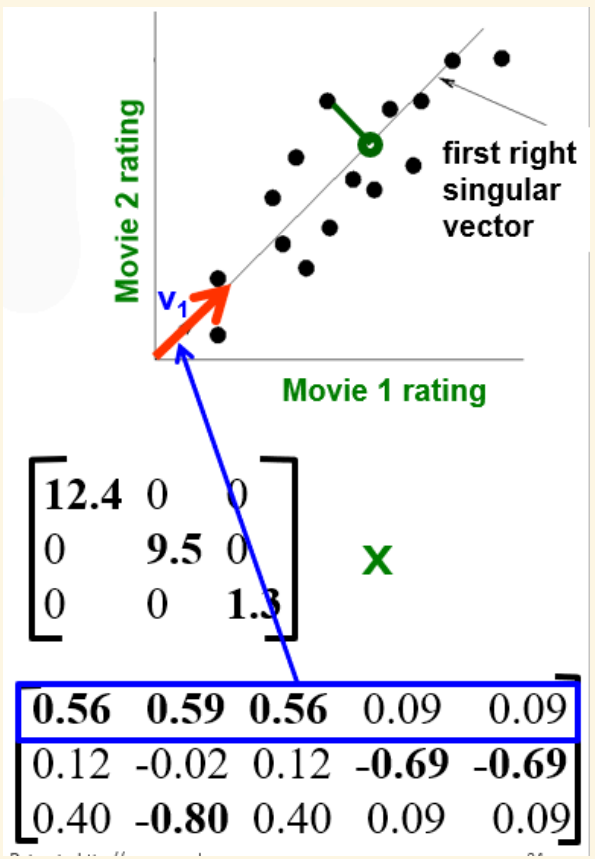

Let's see what happens to the original datapoints

UD project users onto the *concept* axes

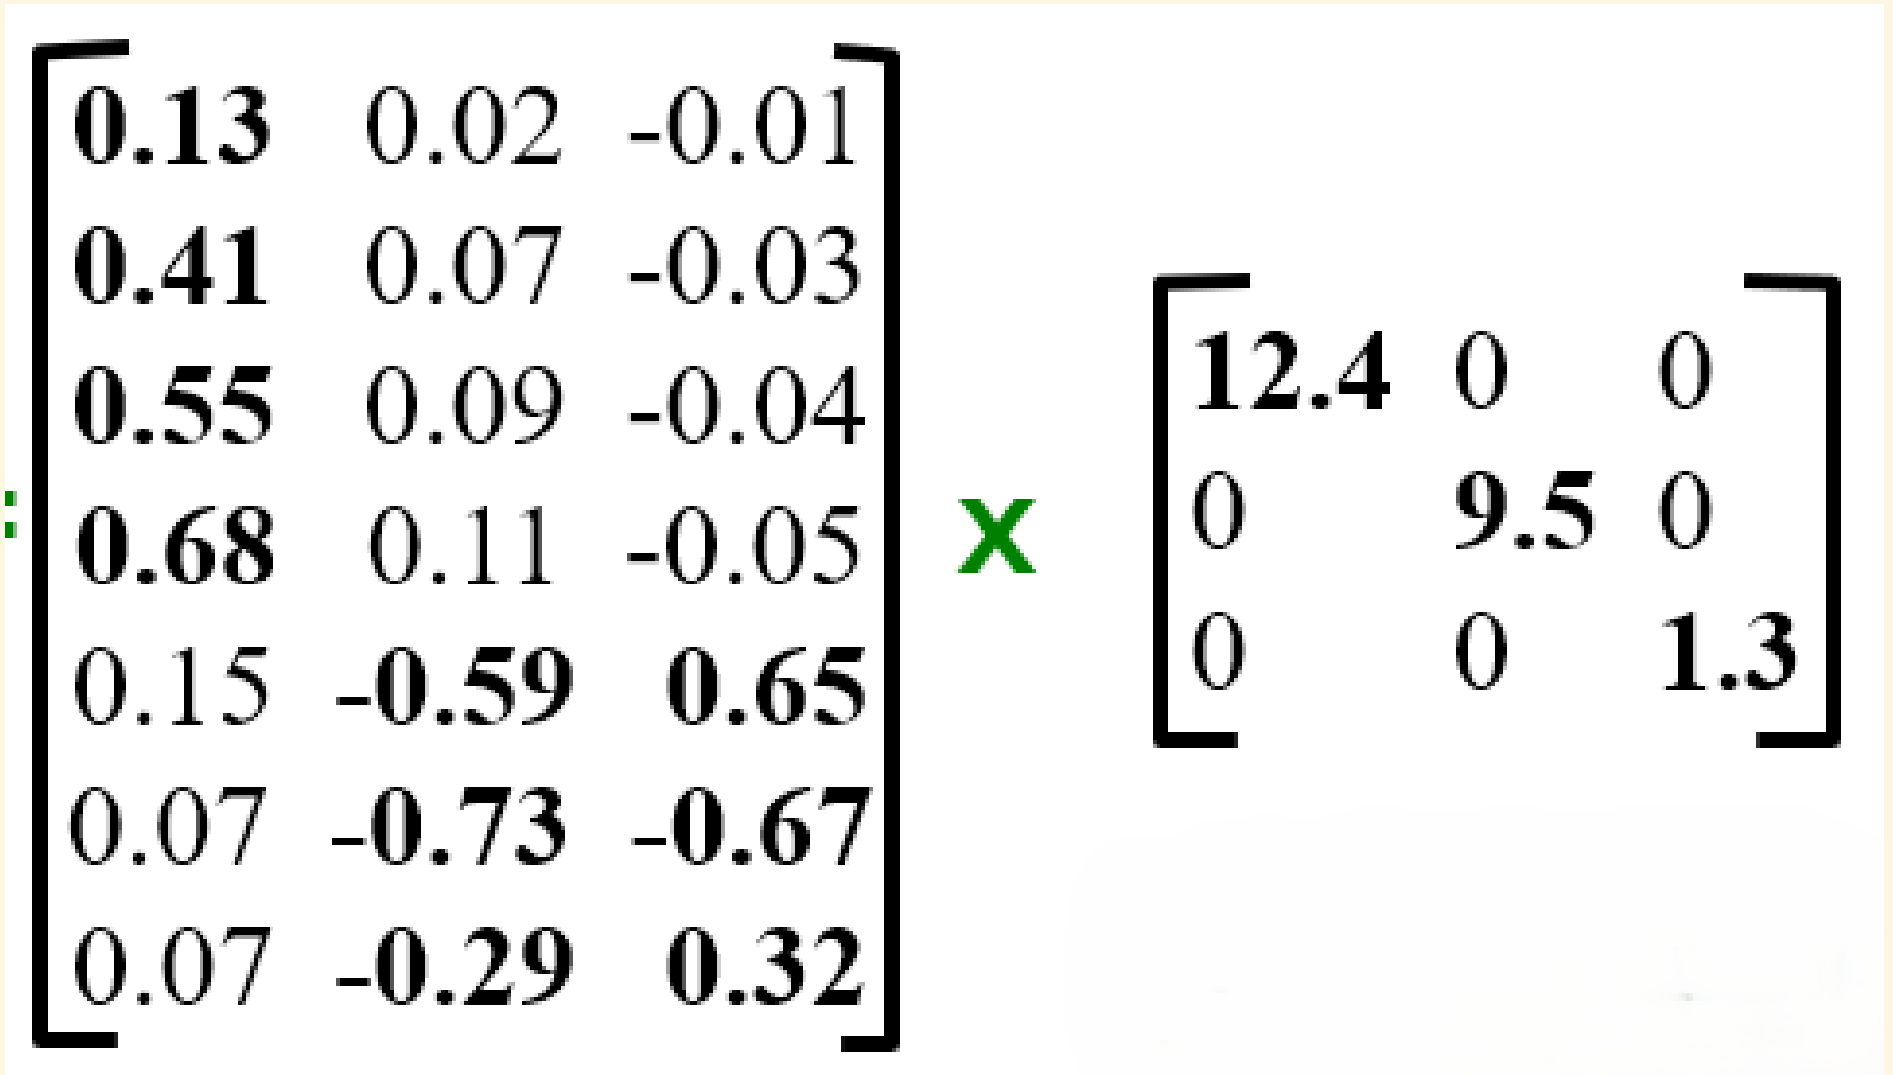

We see values that *look like* measures of the implicit user-to-topic affiliation

Interpreting negative values is problematic, expecially for V^T

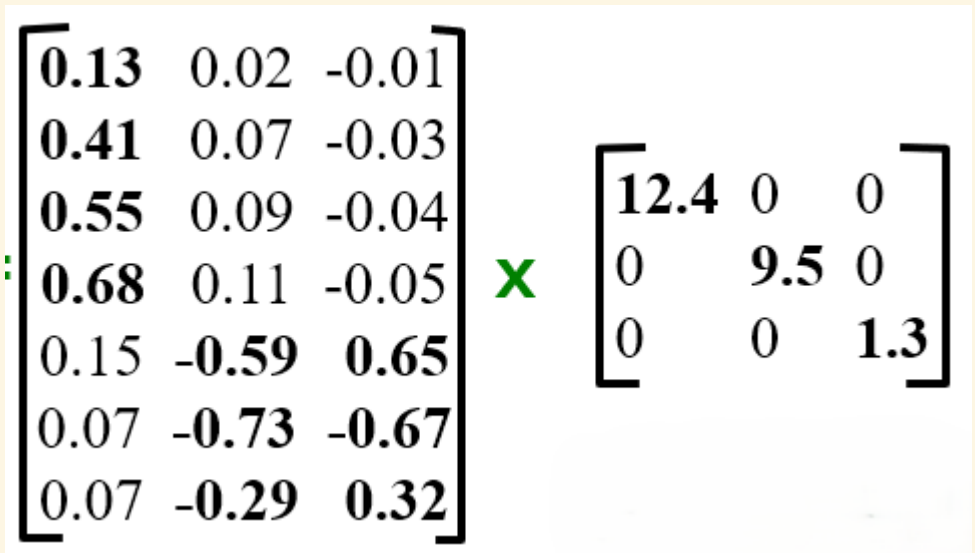

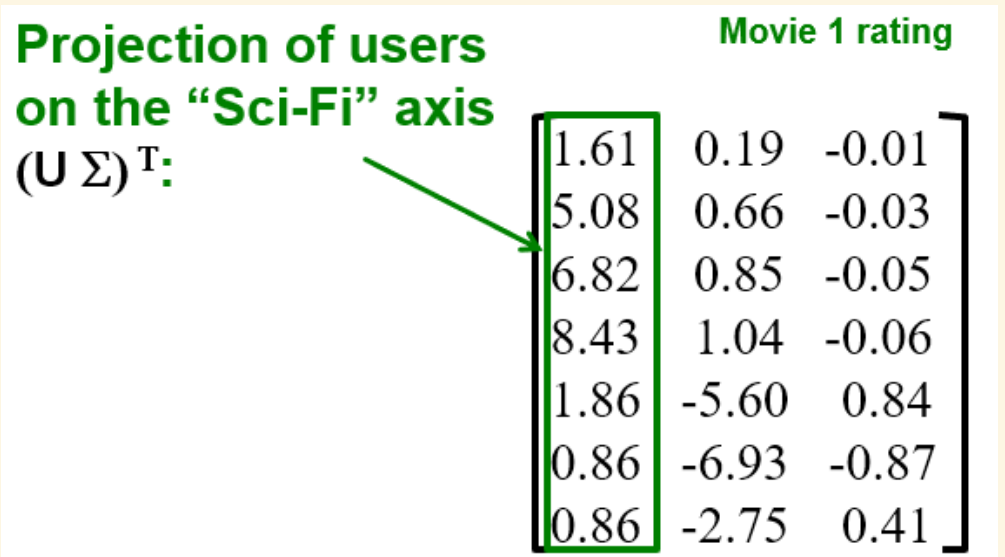

DIMENSIONALITY REDUCTION

SET THE SMALLEST SINGULAR VALUES TO 0

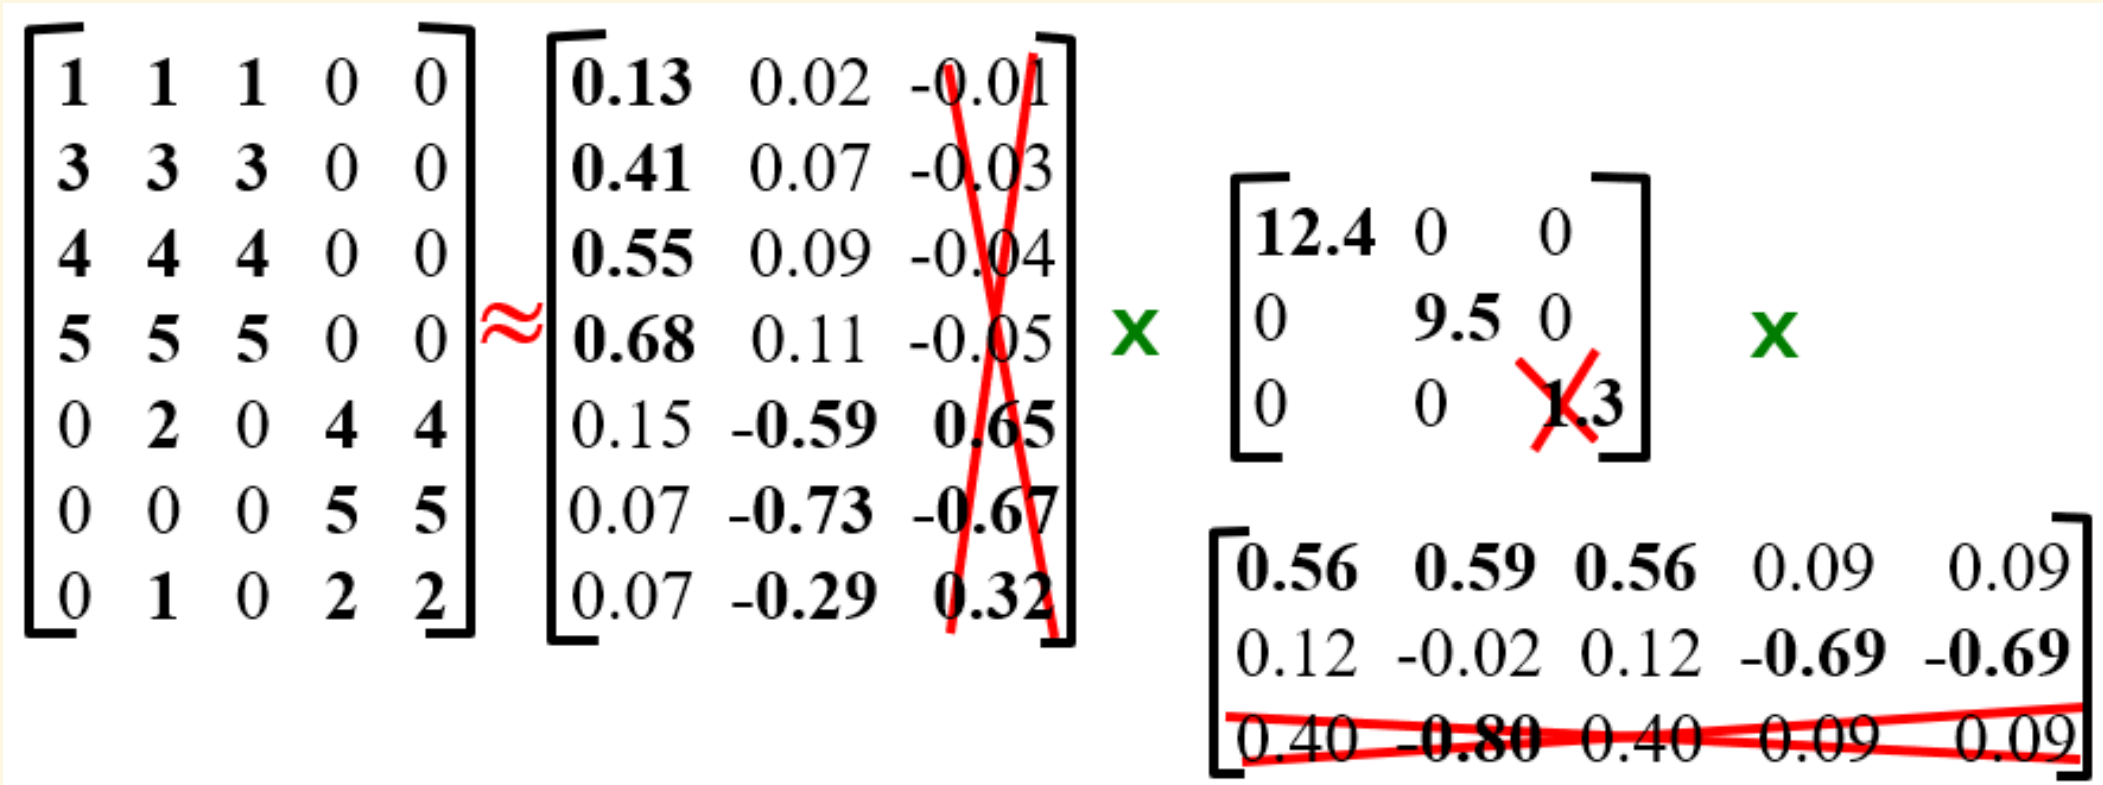

THE B MATRIX

Let $A = UDV^T$ as before, with a ranking $r \leq \min\{m, n\}$

For a small integer k, define

and let E be a *reduction* of D: same everywhere but for $\sigma_{k+1} \ldots \sigma_r$ set to 0.

Theorem: $B = UEV^T$ with ranking k is the best approximation of A with rank k:

$$B \in \operatorname{argmin}_{X:\operatorname{rank}(X)=k}\{||A - X||_F\}$$

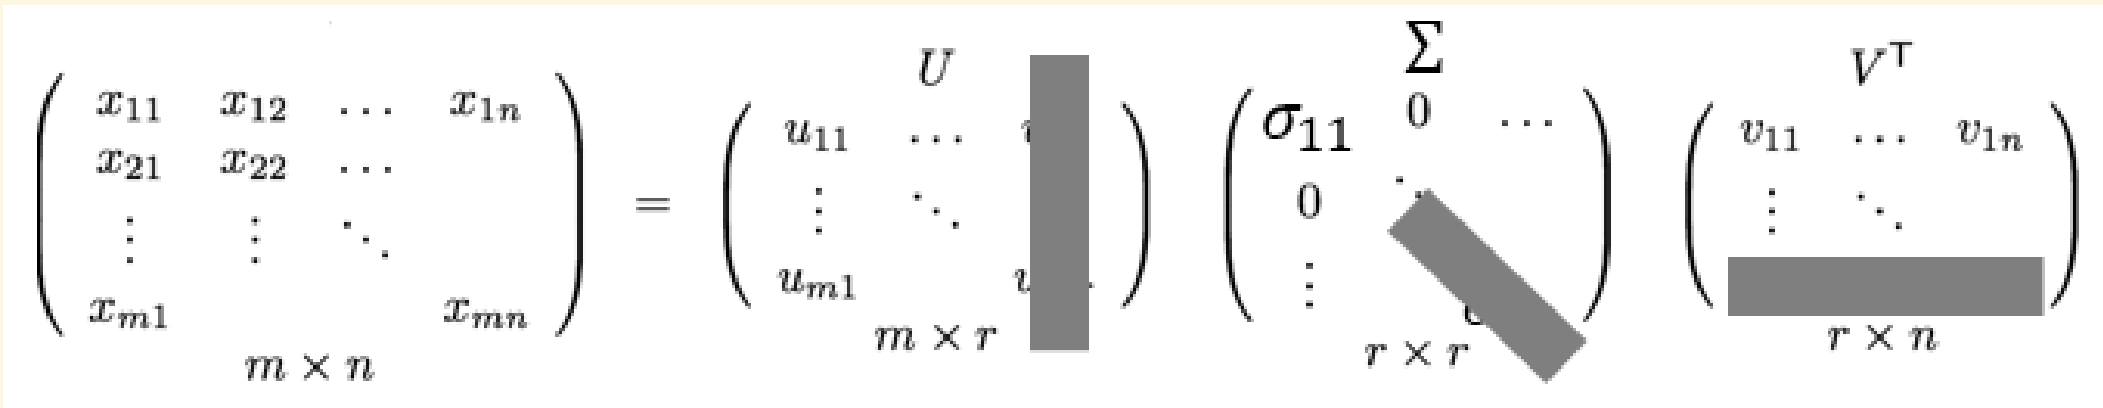

WHY B IS A GOOD REDUCTION?

rows of U and V are the new axes, but they are unit vectors

the singular values σ_i do the scaling

$$A_{(m \times n)} \approx \sigma_1 U_1 \circ V_1^T + \sigma_1 U_1 \circ V_1^T + \ldots \sigma_r U_r \circ V_r^T$$

since we constructed D to have $\sigma_1 \geq \sigma_2 \ldots$, the smallest σ will do little scaling

so dropping $\sigma_{k+1} \ldots \sigma_r$ will introduce less error

Rule of thumb: keep singular values that sum up to about 80% of the total *'energy'*

$$\sum_{1}^{k} \sigma_i \approx 0.8 \sum_{1}^{r} \sigma_j$$

CONCLUSIONS

$A = UDV^T$ provides a unique decomposition that is *interpretable*

Interpretability:

U: user-to-concept affiliations/similarities

D: strenght of each concept

V: film-to-concept affiliations/similarities

SVD picks up linear dependencies

Dimensionality reduction:

SVD finds the *best* reduced matrix B

k is not really an hyperparameter (as with K-nn): we set it empirically to keep 80% of values

SVD dim. reduction is the key to *denoising*

Cost:

the core component requires $\min\{nm^2, n^2m\}$ ops.

Implementation:

`truncatedSVD` in Scikit-learn

```
from sklearn.decomposition import TruncatedSVD

K=2

% X is assigned ...

svd = TruncatedSVD(n_components=K)

X_reduced = svd.fit_transform(X)
```